

Quantization of ecological interactions yields insights into ecosystem assembly and dynamics

Justin D. Yeakel^{a,b}, Mathias Pires^c, James O'Donnell^d, Marcus de Aguiar^c, Paulo R. Guimarães Jr.^e, Dominique Gravel^f, and Thilo Gross^g

^aSchool of Natural Sciences, University of California Merced, Merced, CA 95343, USA; ^bSanta Fe Institute; ^cUniversidade Estadual de Campinas; ^dUniversity of Washington; ^eUniversidade de São Paulo; ^fUniversité de Sherbrooke; ^gUniversity of California Davis, Davis CA

This manuscript was compiled on February 13, 2019

In this paper we investigate the effects of heterogeneity in food available on the survival times of foraging individuals, using a simple model to incorporate the basic elements of time-stationary foraging process.

Ecosystem assembly | Ecological networks | Colonization | Extinction | Ecosystem engineering

Simplifying the abundant complexity and idiosyncrasies of nature is necessary to unravel their secrets. The layers of natural history giving rise to an ecological community can be simplified to a network, where nodes represent species and links represent interactions between them. This minimalist perspective provides insight into community structure, and consequently, the dynamics of species' populations and the stability of the system as a whole (refs). Community structure can directly impact the ability of the system to absorb external perturbations (? ?), the existence and ubiquity of tipping points that can cause sudden long-term changes in species' populations (? ?), and promote or inhibit extinction cascades (?). Importantly, community structure is the result of a dynamical process that takes place over time, where species are added to or removed from the community in succession. This process is generally referred to as community assembly (ref), and though it is an engine for much that we observe in nature, is not well understood.

Ecological networks are generally constructed and analyzed with respect to a single type of interaction. For example, mutualisms form the foundation of mutualistic networks, whereas trophic interactions form the foundation of food webs. More recently, there has been a growing interest in understanding community structure by taking into account diverse interaction types, where multiple unique interactions are included in a multilayer network (?). However, interactions between two species are inherently compound in nature, such that both species in an interaction experiences different effects. For example, in a mutualistic interaction, there is often biomass flow in one direction and a reproductive service in the other; in a trophic interaction there is biomass flow from the prey to the predator and this effect is asymmetric. While this asymmetry is generally not encoded in the structure of ecological networks, it becomes central with description

of network dynamics (May, Allesina, Gross, etc).

Diverse interactions occur directly between species, but also indirectly through the effect that species have on their environment. If these environmental effects have timescales longer than that of the instigating species, we call those species ecosystem engineers (ref). For example, burrowing rodents both create shelter for non-burrowing organisms and aerate the soil, promoting the growth of primary production. Terrestrial plants have altered the stoichiometry of the atmosphere, providing the molecular stores required for respiration (refs). Elephants root out large saplings and small trees, enabling the formation and maintenance of grasslands (refs) on which grass specialists depend. These example describe indirect interactions between ecosystem engineers and species that rely on their engineering, through an environmental intermediary. By way of their indirectness, these interactions are difficult to document. And while the importance of ecosystem engineers is celebrated, explorations of their effects on community dynamics occupy a very limited role in theory (Hastings refs), and are absent in models of ecological networks.

Here we model an ecological system with a network where nodes represent ecological entities, including non-engineering species, engineering species, and the effects of the latter on the environment, which we call *objects*. The links of the network that connect both species and objects represent resource, service, and engineering dependencies, respectively. Following Pillai (REF), we do not track the abundances of entities but only track their presence or absence. We use this framework to explore the dynamics of ecosystem assembly, where the colonization and extinction of species within a community depends on the constraints imposed by the resource, service, and

Significance Statement

Here

Please provide details of author contributions here.

Please declare any conflict of interest here.

²To whom correspondence should be addressed. E-mail: jyeakel@ucmerced.edu

engineering dependencies connecting both species and objects.

We show that our framework reproduces many features of empirical systems both over the course of assembly and at steady state. First, (assembly) proportion generalists Second, (assembly) trophic levels and degree distributions Third, (SS) limited trophic levels, Fourth (SS) increasing mutualisms increases nestedness.

Because we include ecosystem engineering explicitly, we can examine directly the effects of engineers on the assembly of communities. Ecosystem engineers play important roles in assembly processes such as succession, and the maintenance of biodiversity (refs). However, explorations of their effects on community dynamics occupy a very limited role in theory (Hastings refs), and are absent in models of ecological networks. Our framework provides three key insights on the role of ecosystem engineers on ecological communities. First, Second, Third.

Assembly of the ENigMa network

A combination of two directional interactions between a pair of species defines the ecological relationship linking both. A mutualistic interaction between plant x and pollinator y can thus be deconstructed into an ‘eat’ interaction from species x to y and a *need* interaction from species y to x . The former represents a directional trophic interaction, or a resource dependency, whereas the latter represents a reproductive service, or a dependency that does not involve biomass flow. Another type of mutualism is one defined by non-resource dependencies for both species, often called a service-service mutualism. In this case, each species would be connected to the other by directional *need* interactions. A trophic interaction between species x and y can be deconstructed by an *eat* interaction from the consumer x to the prey y , but the absence of an interaction from prey y to the consumer x , as the prey does not interact meaningfully with the consumer. Some species pairs may participate in bi-directional trophic interactions, where the consumer/resource role changes over the course of each species’ life history. In this case, each species would be connected to the other by directional *eat* interactions.

Ecosystem engineering can be formalized by the creation of an object by an engineering species that can be interacted with by other members of a community. Objects made by ecosystem engineers represent any alteration made to the environment, whether this be changes to the stoichiometry of the atmosphere (e.g. primary producers; refs), physical modifications to the landscape (e.g. beavers, salmon; refs), or the production of unique habitats for other organisms (e.g. tree trunks; refs). Objects can be represented in ecological network alongside species as nodes, but where they must be connected to engineering species by a *make* interaction. The *make* interaction between an engineering species and the engineered object denotes the dependency of the environmental effect on the effector.

Diverse interactions without engineers

Community assembly in the absence of engineers reveals the emergence of food web and mutualistic network properties consistent with observations of assembling and steady state empirical systems. Because only primary producers that do not have outgoing need interactions can colonize initially, a diverse base of autotrophic species typically constitutes the early assembly process. In order for communities to have > 1 pure autotroph, we do not consider competitive exclusion of the basal resource, such that all non-mutualistic pure autotrophs have the potential to co-exist. Following the establishment of a suite of autotrophs, both mixotrophs and lower trophic-level heterotrophs assemble into the community (Fig. ??a). As species richness increases, available resources accumulate consumers and competitive exclusion leads to an increase in the extinction rate until a steady state is reached at $S_A^* = 130$ species. This community steady state increases as the number of mutualisms established in the source pool decreases (lower $E\{f_n\}$) because mutualisms introduce dependencies that inhibit colonization.

As the community assembles, we find that the connectance of trophic interactions ($C = L/S^2$) where in our case $L = \sum_{i,j} e_{A(i,j)}$ and $S = S_A$) follows a decay-like trajectory to values similar to – but on average 9% greater than – the connectance of the source pool P (Fig. ??b). Decaying connectance has been documented in the assembly of mangrove communities (Piechnik), however this decay is a statistical inevitability, as a growing food web early in the assembly process inevitably has high link density (few species that are nearly fully connected), and over the course of increasing species richness and the establishment of different trophic levels and compartments, a decline in link density. That the connectance of assembled communities is greater than the source pool is due to the fact that only species connected by trophic interactions can enter the community to begin with, increasing expected link density compared to the overall pool.

Recent empirical work has indicated that generalist species may play an important role early in community assembly, whereas specialists tend to colonize after a diverse resource base has accumulated. Because the definition of a specialist or generalist to some degree depends on the size and connectance of the larger food web, trophic generality can be defined as $G_i = \sum_j e_{A(i,j)}(L/S)^{-1}$, such that the number of trophic interactions for a consumer is scaled by the average number of trophic interactions per species in the community L/S (Piechnik, others). A species is classified as a generalist if the number of its trophic interactions is greater than the average number of links per species, or $G_i > 1$, and a specialist if $G_i < 1$, where a community can be described by the proportion of specialists found therein.

For interaction networks that are assembling over time, generality can be scaled by a number of different measures of L/S , and this has a large effect on our interpretation of the role of generality in community assembly. For

instance, L/S may be quantified by either including all autotrophic species or only autotrophic functional groups. Furthermore, the scaling of generality may be made with respect to the current state of the community at each point in time, or with respect to the community at steady state. For instance, in their investigation of assembling mangrove food webs (originally described by Simberloff, xxx), Piechnik et al. (2008) scaled trophic breadth to a standard steady state value of $L^*/S^* = 0.2$ averaged across 102 food webs. To examine how our assessment of the role of generalism over the course of assembly changes based on the application of different scalings, we employ three different measures of L/S to calculate G_i : 1) G_i^{all} , where L accounts for all links in the food web and S accounts for all species relative to each time interval in the assembly process (circles; Fig. 3b); 2) G_i^{hetero} , where we consider only the links and species richness of heterotrophs, excluding autotrophs (points; Fig. 3b); 3) G_i^* , where L and S are measured with respect to the communities at steady state, which is most similar to the measure used to evaluate assembling mangrove food webs (diamonds; Fig. 3b).

Whether trophic breadth is scaled to the current state of L/S or the steady state value of L^*/S^* has a large influence on the estimated proportion of generalists in the community, particularly when the size of the system is small. We observe that for G_i^{all} , the system is initially assembled by specialist species, though over the course of assembly the proportion of specialists relative to generalists declines to intermediate values (circles representing the average over replicates in Fig. 3). If only the trophic links between non-autotrophs are considered as in G_i^{hetero} , specialists still dominate early in assembly, but there is a greater range, such that some systems can be described by a mixed proportion of specialists and generalists (individual points representing independent replicates in Fig. 3). If generalism is measured with respect to the steady state L^*/S^* as in G_i^* , we observe that generalists dominate early in assembly, with an increase in specialists as assembly progresses (diamonds representing the average over replicates in Fig. 3). At steady state, all measures of L/S are approximately equivalent, and the proportion of specialists levels out at ca. 60%, which is similar to the empirical observations for Simberloff’s mangrove communities in Piechnik et al. (2008).

The role of specialists early in assembly is primarily due to the accumulation of autotrophs specializing on the basal resource. This is evident when we observe that the trophic level distribution early in assembly is peaked at the lowest trophic level (trophic level 1). Four trophic levels are typically established by $t = 50$, where colonization is still the dominant dynamic, and by the time communities reach steady state the interaction networks are characterized by ca. 10 trophic levels (Fig. 1). The distribution of trophic levels changes shape over the course of assembly: early on, we observe that the community exhibits a pyramidal structure, where the vast majority of




Fig. 1. The frequency distribution of trophic levels over the course of assembly taken across 10^4 replicates. Autotrophs occupy trophic level 1. We set $S = 200$, $p_e = 0.01$, $p_n = 0.002$, and $t = 4000$ (see methods).

species inhabit low-trophic positions. At steady state, we observe that intermediate trophic levels (2-7) dominate, with frequencies that reveal an hour-glass structure. We emphasize that these structures are diversity-weighted rather than biomass or abundance-weighted as is often the case (Trebilco et al. 2013, Gibert & Yeakel 2019). Trophic levels higher than 7 do occur, but are increasingly rare.

Because the ENIgMa framework includes multiple types of interactions, we must examine whether structures characteristic of mutualistic networks are observed. Empirical observations of mutualistic systems reveal that such interactions tend to be nested (where specialist interactions are subsets of generalist interactions). Increasing the frequency of need interactions increases the frequency of both service-resource ($e \leftrightarrow n$) and service-service ($n \leftrightarrow n$) mutualisms. In the ENIgMa framework, the probability of competitive exclusion is reduced with an increase in the prevalence mutualistic interactions, and this should lead more stable nested bipartite motifs over the course of assembly (see Supplementary Information). Our expectation then, is that nestedness should increase with the frequency of mutualisms, though this is difficult to predict a priori. As we increase the frequency of need interactions in the source pool, we indeed observe an increase in nestedness (measured as NODF; Fig. 2). That the absolute values of nestedness are low compared to those measured for empirical mutualistic networks is unsurprising: observations of mutualistic interactions are generally for bipartite networks and isolated to specific systems (e.g. ant-plant mutualisms). Here the NODF metric is taken across both eat and need interactions across the entire assembled community.

The role of engineers




Fig. 2. Nestedness (measured as NODF) as a function of the frequency of need interactions in the species pool interaction matrix P .

Beyond engineering

Materials and Methods

We model an ecological system with a network where nodes represent *ecological entities* such as populations of species and or the presence of inanimate objects affecting species such as (examples). Following Pillai (REF), we do not track the abundances of entities but only track their presence or absence. The links of the network represent interactions between pairs of entities (x,y) . We distinguish three types of such interactions: x eats y , x needs y to be present, x makes object y .

The model is initialized by creating S species and $O = \eta S$ objects, such that $N = S + O$ is the total number of entities and η is the number of objects per species in the system. For each pair of species (x,y) there is a probability p_e that x eats y and probability p_n that x needs y . For each pair of species x and object o , there is a probability q_e that x eats o and a probability q_n that species x needs object o . Additionally, each species makes a number of objects that is drawn from a Poisson distribution with mean $\mu = \eta e(e-1)^{-1}$ where e is Euler's number. Once the number of objects per species is determined, each object is assigned to a species independently. This means that there multiple species may make the same object, and that may be some objects that are not made by any species.

In addition to interactions with ecosystem entities, there can be interactions with a basal resource, which is always present. The first species always eats this resource, such that there is always a primary producer in the pool. Other species eat the basal resource with probability p_e .

We then consider the assembly of a community which at any time will contain a subset of entities in the pool and always the basal resource. In time, the entities in the community are updated following a set of rules. A species from the pool can colonize the community if the following conditions are met: 1) all entities that a species needs are present in the community, and 2) at least one entity that a species eats is present in the community. If a colonization event is possible, it occurs stochastically in time with rate r_c .

An established species is at risk of extinction if it is not the strongest competitor at least one of its resources that it eats. We compute the competitive strength of species i as

$$\sigma_i = c_n n_i - c_e e_i - c_v v_i, \quad [1]$$

where n_i is the number of entities that species i needs, e_i is the number of entities from the pool that species i can eat, and v_i is the number of species in the community that eat species i .




Fig. 3. The proportion of specialists as a function of assembly time, where a specialist is defined as a species with a generality index $G_i < 1$. All measures of G_i are scaled by the average number of links per species L/S , and we consider different values of L/S on G_i : Circles: G_i^{all} where L accounts for all links in the food web and S accounts for all species relative to each time interval in the assembly process (averaged across replicates); Points: G_i^{hetero} , where we consider only the links and species richness of heterotrophs, excluding autotrophs (each point shows an individual replicate); Diamonds: G_i^* , where L and S are measured with respect to the communities at steady state, which is most similar to the measure used to evaluate assembling mangrove food webs (averaged across replicates).

This captures the ecological intuition that mutualisms provide a fitness benefit, specialists are stronger competitors than generalists, and many predators entail an energetic cost. The coefficients c_n , c_e , c_v describe the relative effects of these contributions to competitive strength. In the following, we use the values $c_n = \pi$, $c_e = \sqrt{2}$, $c_v = 1$, such that the competitive benefit of adding an additional mutualism is greater than the detriment incurred by adding another prey or predator. A species at risk of extinction leaves the community stochastically in time at rate r_e .

An object is present in the community whenever at least one species that makes the object is present. If a species that makes an object colonizes a community, the object is created immediately, however objects may persist for some time after the last species that makes the object goes extinct. Any object that has lost all of its makers disappears stochastically in time at rate r_o .

The model described here can be simulated efficiently with an event-driven simulation utilizing a Gillespie algorithm. In these types of simulations, one computes the rates r_j of all possible events j in a given step. One then selects the time at which the next event happens by drawing a random number from an exponential distribution with mean $1/\sum_j r_j$. At this time, an event occurs that is randomly selected from the set of possible events such that the probability of event a is $r_a/\sum_j r_j$. Then the effect of the event is realized and the list of possible events is updated for the next step. This algorithm is known to offer a much better approximation to the true stochastic continuous time process than a simulation in discrete time steps, while providing a much higher numerical efficiency (Gillespie, 1974ish).

Parameter	Definition	Value/Range
\vec{a}	assimilate	
\vec{n}	need	
\vec{i}	ignore	
\vec{m}	make	
$e \leftrightarrow i$	Asymmetric consumption	$p_{ei} = p_i(p_e/(p_e + p_n + p_i)) + p_e(p_i/(p_a + p_i + p_n))$
$e \leftrightarrow e$	Symmetric consumption	$p_{ee} = p_e(p_e/(p_i + p_n + p_e))$
$e \leftrightarrow n$	Trophic mutualism	$p_{en} = p_n(p_e/(p_e + p_n + p_i + p_m)) + p_e(p_n/(p_a + p_i + p_n))$
$n \leftrightarrow n$	Non-trophic mutualism	$p_{nn} = p_n(p_n/(p_e + p_n + p_i + p_m))$
$n \leftrightarrow i$	Commensalism	$p_{ni} = p_n(p_i/(p_e + p_n + p_i + p_m)) + p_i(p_n/(p_e + p_n + p_i))$
$m \leftrightarrow n$	Engineering	$p_{mn} = p_n(p_m/(p_e + p_n + p_i + p_m)) + p_m$
$i \leftrightarrow i$	Null	$p_{ii} = p_i(p_i/(p_e + p_n + p_i))$
\mathcal{N}	Number of species + objects	dyn.
\mathcal{S}	Number of species	dyn.
\mathcal{O}	Number of objects	dyn.

Table 1. Table of parameters, definitions, and assigned values or ranges.




Fig. 4. An example of the source pool interaction matrix where $\mathcal{S} = 50$. Species and objects are aligned across the rows and columns; objects are shaded and labeled by 'o' to distinguish them from species. The interaction recorded in row i and column j describes the directed interaction from species/object i to species/object j . The first row/column represents the basal resource; species that assimilate the primary resource are capable of primary production. Species interact with other species and/or objects; objects only interact with their engineers by 'needing' them; objects do not interact with other objects.